




Mean total tokens



Mean request peak KV cache



Mean client CPU (64 logical CPUs)



Mean latency

